

Advanced Data Analysis

Homework Week 7

Aswin Vijay

June 12, 2023

1] Math

We are asked to derive $\frac{\partial J}{\partial x_i}$ where x_i is the i 'th units pre-activation value and J is the loss. We are given $\frac{\partial J}{\partial u_i}$, where u_i is the output after the batch normalization operation. We also have,

$$\mu_{\mathcal{B}} = \frac{1}{m} \sum_{i=1}^m x_i \quad \text{mini - batch mean} \quad (1)$$

$$\sigma_{\mathcal{B}}^2 = \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad \text{mini - batch variance} \quad (2)$$

$$\hat{x}_i = \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad \text{normalized } x_i \quad (3)$$

$$u_i = \gamma \hat{x}_i + \beta \quad \text{Scale and shift} \quad (4)$$

To derive $\frac{\partial J}{\partial x_i}$ we write out the partial derivatives using chain rule since $\hat{x}_i(x_i, \mu, \sigma^2)$, $\sigma^2(x_i, \mu)$ and $\mu(x_i)$, three variables are dependent on x_i so:

$$\frac{\partial J}{\partial x_i} = \frac{\partial J}{\partial \hat{x}_i} \cdot \frac{\partial \hat{x}_i}{\partial x_i} + \frac{\partial J}{\partial \mu} \cdot \frac{\partial \mu}{\partial x_i} + \frac{\partial J}{\partial \sigma^2} \cdot \frac{\partial \sigma^2}{\partial x_i}$$

The above derivatives are computed as follows,

$$\begin{aligned} \frac{\partial J}{\partial \hat{x}_i} &= \frac{\partial J}{\partial u_i} \cdot \frac{\partial u_i}{\partial \hat{x}_i} = \gamma \cdot \frac{\partial J}{\partial u_i} \quad \text{From (4)} \\ \frac{\partial \hat{x}_i}{\partial x_i} &= \frac{1}{\sqrt{\sigma^2 + \epsilon}} \quad \text{From (3)} \end{aligned}$$

Computing J derivatives w.r.t μ ,

$$\begin{aligned}\frac{\partial J}{\partial \mu} &= \sum_{i=1}^m \frac{\partial J}{\partial \hat{x}_i} \cdot \frac{\partial \hat{x}_i}{\partial \mu} + \frac{\partial J}{\partial \sigma^2} \cdot \frac{\partial \sigma^2}{\partial \mu} \text{ Summing over batch} \\ \frac{\partial \hat{x}_i}{\partial \mu} &= \frac{-1}{\sqrt{\sigma^2 + \epsilon}} \text{ From (3)} \\ \frac{\partial \sigma^2}{\partial \mu} &= \frac{-2}{m} \sum_{i=1}^m (x_i - \mu) = -2 \left(\frac{1}{m} \sum_{i=1}^m x_i - \mu \right) = 0 \text{ From (2,1)}\end{aligned}$$

Computing J derivatives w.r.t σ ,

$$\begin{aligned}\frac{\partial J}{\partial \sigma^2} &= \sum_{i=1}^m \frac{\partial J}{\partial \hat{x}_i} \cdot \frac{\partial \hat{x}_i}{\partial \sigma^2} \text{ Summing over batch} \\ \frac{\partial \hat{x}_i}{\partial \sigma^2} &= \frac{-1}{2} (\sigma^2 + \epsilon)^{-\frac{3}{2}} (x_i - \mu) \text{ From (3)}\end{aligned}$$

Putting everything together we have,

$$\begin{aligned}\frac{\partial J}{\partial \mu} &= \sum_{i=1}^m \frac{\partial J}{\partial \hat{x}_i} \cdot \frac{-1}{\sqrt{\sigma^2 + \epsilon}} + 0, \text{ then} \\ \frac{\partial J}{\partial x_i} &= \frac{\partial J}{\partial \hat{x}_i} \cdot \frac{1}{\sqrt{\sigma^2 + \epsilon}} + \sum_{i=1}^m \frac{\partial J}{\partial \hat{x}_i} \cdot \frac{-1}{\sqrt{\sigma^2 + \epsilon}} \cdot \frac{\partial \mu}{\partial x_i} + \frac{-1}{2} (\sigma^2 + \epsilon)^{-\frac{3}{2}} \sum_{j=1}^m \frac{\partial J}{\partial \hat{x}_j} \cdot (x_j - \mu) \cdot \frac{\partial \sigma^2}{\partial x_i}\end{aligned}$$

The remaining derivatives are calculated as below,

$$\begin{aligned}\frac{\partial \mu}{\partial x_i} &= \frac{1}{m} \text{ From (1)} \\ \frac{\partial \sigma^2}{\partial x_i} &= \frac{2(x_i - \mu)}{m} \text{ From (2)}\end{aligned}$$

Finally we get,

$$\begin{aligned}\frac{\partial J}{\partial x_i} &= \frac{\partial J}{\partial \hat{x}_i} \cdot \frac{1}{\sqrt{\sigma^2 + \epsilon}} + \frac{1}{m} \sum_{i=1}^m \frac{\partial J}{\partial \hat{x}_i} \cdot \frac{-1}{\sqrt{\sigma^2 + \epsilon}} - \frac{(\sigma^2 + \epsilon)^{-\frac{3}{2}}}{m} \sum_{j=1}^m \frac{\partial J}{\partial \hat{x}_j} \cdot (x_j - \mu) \cdot (x_i - \mu) \\ &= \frac{\partial J}{\partial \hat{x}_i} \cdot \frac{1}{\sqrt{\sigma^2 + \epsilon}} + \frac{1}{m} \sum_{i=1}^m \frac{\partial J}{\partial \hat{x}_i} \cdot \frac{-1}{\sqrt{\sigma^2 + \epsilon}} - \frac{1}{m\sqrt{\sigma^2 + \epsilon}} \cdot \hat{x}_i \cdot \sum_{j=1}^m \frac{\partial J}{\partial \hat{x}_j} \hat{x}_j \\ &= \frac{1}{m\sqrt{\sigma^2 + \epsilon}} \left(m \frac{\partial J}{\partial \hat{x}_i} - \sum_{i=1}^m \frac{\partial J}{\partial \hat{x}_i} - \hat{x}_i \sum_{j=1}^m \frac{\partial J}{\partial \hat{x}_j} \hat{x}_j \right)\end{aligned}$$

2] Architecture

In Convolutional Neural Networks the bias parameter is associated with each filter and its value is added to the filter output. Such a bias is called a tied

bias where the bias remains constant for each location of a feature map, here the learnable parameters is lesser. Untied biases can also be used in which case each location on the input map can have its own bias, the number of learnable parameters drastically increases but the network would now be able to learn location specific information allowing for more fine tuning. Their use depends on the training data, if the data is shifted uniformly over all features a tied bias would work well, but if it has location specific shifts an untied bias can be used for correction.

If there is batch normalization operation after the convolution then biasing would be point less as the bias would also get normalized in the case of tied biases. The bias term thus becomes redundant. In case of untied biases the effect will remain as each neuron gets biased differently.

ADA7-ThreeLayer.py

```
1  # Implementation of a Three-layer neural network for MNIST classification
2  # with manual gradient calculation and manual optimization.
3
4  import torch
5  import torch.nn.functional as F
6  from torch.utils.data import DataLoader
7  from torchvision import datasets, transforms
8  import matplotlib.pyplot as plt
9
10 torch.manual_seed(0)
11 lr = 0.005
12 hidden_dim = 15
13 batch_size = 64
14 epochs = 20
15 plot = True
16
17 transform = transforms.Compose([transforms.ToTensor(), transforms.Normalize((0.1307,),
18 (0.3081,))]) # mnist mean and std
19 train_dataset = datasets.MNIST(root='./data', train=True, download=True,
20 transform=transform)
21 test_dataset = datasets.MNIST(root='./data', train=False, download=True,
22 transform=transform)
23 train_loader = DataLoader(dataset=train_dataset, batch_size=batch_size, shuffle=True)
24 test_loader = DataLoader(dataset=test_dataset, batch_size=1000, shuffle=False)
25
26 class ThreeLayerNet:
27     def __init__(self, input_size, hidden_size1, hidden_size2, output_size):
28         self.W1 = torch.randn(input_size, hidden_size1) * 0.1
29         self.b1 = torch.randn(hidden_size1) * 0.1
30         self.W2 = torch.randn(hidden_size1, hidden_size2) * 0.1
31         self.b2 = torch.randn(hidden_size2) * 0.1
32         self.W3 = torch.randn(hidden_size2, output_size) * 0.1
33         self.b3 = torch.randn(output_size) * 0.1
34
35     def forward(self, x):
36         # Input
37         self.x = x
38         # Hidden Layer 1
39         self.z1 = x @ self.W1 + self.b1
40         self.a1 = F.relu(self.z1)
41         # Hidden Layer 2
42         self.z2 = self.a1 @ self.W2 + self.b2
43         self.a2 = F.relu(self.z2)
44         # Output Layer
45         self.z3 = self.a2 @ self.W3 + self.b3
46         return self.z3
47
48 model = ThreeLayerNet(784, hidden_dim, hidden_dim, 10)
49
50 train_loss_values = []
51 test_loss_values = []
52 train_accuracy_values = []
53 test_accuracy_values = []
54
55 def train(epoch):
56     for batch_idx, (data, target) in enumerate(train_loader):
```

```

55     data = data.view(-1, 784) # flatten the input
56
57     # === FORWARD PASS ===
58     output = model.forward(data)
59     log_softmax = F.log_softmax(output, dim=1)
60     loss = - torch.mean(log_softmax[range(len(target)), target]) # Equivalent to
NLLLoss
61
62     # === BACKWARD PASS ===
63     # gradient of the loss w.r.t. output of model
64     grad_z3 = F.softmax(output, dim=1)
65     grad_z3[range(len(target)), target] -= 1
66     grad_z3 /= len(target) # recall that loss is average over batch
67
68     # gradient of the loss w.r.t. the output after the second hidden layer ReLU
69     grad_a2 = grad_z3 @ model.W3.T
70
71     # gradient of the loss w.r.t. the output before the second hidden layer ReLU
72     grad_z2 = grad_a2.clone()
73     grad_z2[model.z2 < 0] = 0
74
75     # gradient of the loss w.r.t. the output after the first hidden layer ReLU
76     grad_a1 = grad_z2 @ model.W2.T
77
78     # gradient of the loss w.r.t. the output before the first hidden layer ReLU
79     grad_z1 = grad_a1.clone()
80     grad_z1[model.z1 < 0] = 0
81
82     # gradient of the loss w.r.t. the model parameters
83     model.W3.grad = model.a2.T @ grad_z3
84     model.b3.grad = grad_z3.sum(axis=0)
85     model.W2.grad = model.a1.T @ grad_z2
86     model.b2.grad = grad_z2.sum(axis=0)
87     model.W1.grad = model.x.T @ grad_z1
88     model.b1.grad = grad_z1.sum(axis=0)
89
90     # === PARAM UPDATES === # Vanilla gradient descent
91     model.W3 -= lr * model.W3.grad
92     model.b3 -= lr * model.b3.grad
93     model.W2 -= lr * model.W2.grad
94     model.b2 -= lr * model.b2.grad
95     model.W1 -= lr * model.W1.grad
96     model.b1 -= lr * model.b1.grad
97
98     # === PRINT EVERY 200 ITERATIONS ===
99     if batch_idx % 200 == 0:
100         print('Train Epoch: {} [{}/{} ({:.0f}%)]\tLoss: {:.6f}'.format(
101             epoch, batch_idx * len(data), len(train_loader.dataset),
102             100. * batch_idx / len(train_loader), loss.item()))
103
104 @torch.no_grad()
105 def check(train_or_test, loader):
106     loss, correct = 0, 0
107     for data, target in loader:
108         data = data.view(-1, 784) # flatten the input
109         output = model.forward(data)
110         log_softmax = F.log_softmax(output, dim=1)
111         loss += F.nll_loss(log_softmax, target, reduction='sum').item() # sum up batch
loss
112     pred = output.argmax(dim=1, keepdim=True) # get the index of the max log-
probability

```

```

113     correct += pred.eq(target.view_as(pred)).sum().item() # pred is batch_size x 1,
target is batch_size
114
115     loss /= len(loader.dataset) # note that reduction is 'sum' instead of 'mean' in
F.nll_loss
116     accuracy = 100. * correct / len(loader.dataset)
117     print('{} set: Average loss: {:.4f}, Accuracy: {}/{ } ({:.0f}%)' .format(
118         train_or_test, loss, correct, len(loader.dataset), accuracy))
119     return loss, accuracy
120
121 for epoch in range(1, 1+epochs):
122     train(epoch)
123     train_loss, train_acc = check(train_or_test='train', loader=train_loader)
124     train_loss_values.append(train_loss)
125     train_accuracy_values.append(train_acc)
126     test_loss, test_acc = check(train_or_test='test', loader=test_loader)
127     test_loss_values.append(test_loss)
128     test_accuracy_values.append(test_acc)
129     print('', end='\n')
130     # Notice how we are doing a full forward pass again at the end of each epoch to
compute the train in loss and accuracy,
131     # but to save time, we could keep track of the loss (or number of correct predictions)
for each mini-batch
132     # and take an average at the end of the epoch instead.
133
134 if plot:
135     plt.figure(figsize=(12, 5))
136     plt.subplot(1, 2, 1)
137     plt.plot(range(1, 1+epochs), train_loss_values, label='Training Loss')
138     plt.plot(range(1, 1+epochs), test_loss_values, label='Test Loss')
139     plt.xlabel('Epochs')
140     plt.ylabel('Loss')
141     plt.legend()
142
143     plt.subplot(1, 2, 2)
144     plt.plot(range(1, 1+epochs), train_accuracy_values, label='Training Accuracy')
145     plt.plot(range(1, 1+epochs), test_accuracy_values, label='Test Accuracy')
146     plt.xlabel('Epochs')
147     plt.ylabel('Accuracy (%)')
148     plt.legend()
149
150     plt.show()

```

