

ADA7-ThreeLayer.py

```
1  # Implementation of a Three-layer neural network for MNIST classification
2  # with manual gradient calculation and manual optimization.
3
4  import torch
5  import torch.nn.functional as F
6  from torch.utils.data import DataLoader
7  from torchvision import datasets, transforms
8  import matplotlib.pyplot as plt
9
10 torch.manual_seed(0)
11 lr = 0.005
12 hidden_dim = 15
13 batch_size = 64
14 epochs = 20
15 plot = True
16
17 transform = transforms.Compose([transforms.ToTensor(), transforms.Normalize((0.1307,),
18 (0.3081,))]) # mnist mean and std
19 train_dataset = datasets.MNIST(root='./data', train=True, download=True,
20 transform=transform)
21 test_dataset = datasets.MNIST(root='./data', train=False, download=True,
22 transform=transform)
23 train_loader = DataLoader(dataset=train_dataset, batch_size=batch_size, shuffle=True)
24 test_loader = DataLoader(dataset=test_dataset, batch_size=1000, shuffle=False)
25
26 class ThreeLayerNet:
27     def __init__(self, input_size, hidden_size1, hidden_size2, output_size):
28         self.W1 = torch.randn(input_size, hidden_size1) * 0.1
29         self.b1 = torch.randn(hidden_size1) * 0.1
30         self.W2 = torch.randn(hidden_size1, hidden_size2) * 0.1
31         self.b2 = torch.randn(hidden_size2) * 0.1
32         self.W3 = torch.randn(hidden_size2, output_size) * 0.1
33         self.b3 = torch.randn(output_size) * 0.1
34
35     def forward(self, x):
36         # Input
37         self.x = x
38         # Hidden Layer 1
39         self.z1 = x @ self.W1 + self.b1
40         self.a1 = F.relu(self.z1)
41         # Hidden Layer 2
42         self.z2 = self.a1 @ self.W2 + self.b2
43         self.a2 = F.relu(self.z2)
44         # Output Layer
45         self.z3 = self.a2 @ self.W3 + self.b3
46         return self.z3
47
48 model = ThreeLayerNet(784, hidden_dim, hidden_dim, 10)
49
50 train_loss_values = []
51 test_loss_values = []
52 train_accuracy_values = []
53 test_accuracy_values = []
54
55 def train(epoch):
56     for batch_idx, (data, target) in enumerate(train_loader):
```

```

55     data = data.view(-1, 784) # flatten the input
56
57     # === FORWARD PASS ===
58     output = model.forward(data)
59     log_softmax = F.log_softmax(output, dim=1)
60     loss = - torch.mean(log_softmax[range(len(target)), target]) # Equivalent to
NLLLoss
61
62     # === BACKWARD PASS ===
63     # gradient of the loss w.r.t. output of model
64     grad_z3 = F.softmax(output, dim=1)
65     grad_z3[range(len(target)), target] -= 1
66     grad_z3 /= len(target) # recall that loss is average over batch
67
68     # gradient of the loss w.r.t. the output after the second hidden layer ReLU
69     grad_a2 = grad_z3 @ model.W3.T
70
71     # gradient of the loss w.r.t. the output before the second hidden layer ReLU
72     grad_z2 = grad_a2.clone()
73     grad_z2[model.z2 < 0] = 0
74
75     # gradient of the loss w.r.t. the output after the first hidden layer ReLU
76     grad_a1 = grad_z2 @ model.W2.T
77
78     # gradient of the loss w.r.t. the output before the first hidden layer ReLU
79     grad_z1 = grad_a1.clone()
80     grad_z1[model.z1 < 0] = 0
81
82     # gradient of the loss w.r.t. the model parameters
83     model.W3.grad = model.a2.T @ grad_z3
84     model.b3.grad = grad_z3.sum(axis=0)
85     model.W2.grad = model.a1.T @ grad_z2
86     model.b2.grad = grad_z2.sum(axis=0)
87     model.W1.grad = model.x.T @ grad_z1
88     model.b1.grad = grad_z1.sum(axis=0)
89
90     # === PARAM UPDATES === # Vanilla gradient descent
91     model.W3 -= lr * model.W3.grad
92     model.b3 -= lr * model.b3.grad
93     model.W2 -= lr * model.W2.grad
94     model.b2 -= lr * model.b2.grad
95     model.W1 -= lr * model.W1.grad
96     model.b1 -= lr * model.b1.grad
97
98     # === PRINT EVERY 200 ITERATIONS ===
99     if batch_idx % 200 == 0:
100         print('Train Epoch: {} [{}/{} ({:.0f}%)]\tLoss: {:.6f}'.format(
101             epoch, batch_idx * len(data), len(train_loader.dataset),
102             100. * batch_idx / len(train_loader), loss.item()))
103
104 @torch.no_grad()
105 def check(train_or_test, loader):
106     loss, correct = 0, 0
107     for data, target in loader:
108         data = data.view(-1, 784) # flatten the input
109         output = model.forward(data)
110         log_softmax = F.log_softmax(output, dim=1)
111         loss += F.nll_loss(log_softmax, target, reduction='sum').item() # sum up batch
loss
112     pred = output.argmax(dim=1, keepdim=True) # get the index of the max log-
probability

```

```

113     correct += pred.eq(target.view_as(pred)).sum().item() # pred is batch_size x 1,
target is batch_size
114
115     loss /= len(loader.dataset) # note that reduction is 'sum' instead of 'mean' in
F.nll_loss
116     accuracy = 100. * correct / len(loader.dataset)
117     print('{} set: Average loss: {:.4f}, Accuracy: {}/{ } ({:.0f}%)' .format(
118         train_or_test, loss, correct, len(loader.dataset), accuracy))
119     return loss, accuracy
120
121 for epoch in range(1, 1+epochs):
122     train(epoch)
123     train_loss, train_acc = check(train_or_test='train', loader=train_loader)
124     train_loss_values.append(train_loss)
125     train_accuracy_values.append(train_acc)
126     test_loss, test_acc = check(train_or_test='test', loader=test_loader)
127     test_loss_values.append(test_loss)
128     test_accuracy_values.append(test_acc)
129     print('', end='\n')
130     # Notice how we are doing a full forward pass again at the end of each epoch to
compute the train in loss and accuracy,
131     # but to save time, we could keep track of the loss (or number of correct predictions)
for each mini-batch
132     # and take an average at the end of the epoch instead.
133
134 if plot:
135     plt.figure(figsize=(12, 5))
136     plt.subplot(1, 2, 1)
137     plt.plot(range(1, 1+epochs), train_loss_values, label='Training Loss')
138     plt.plot(range(1, 1+epochs), test_loss_values, label='Test Loss')
139     plt.xlabel('Epochs')
140     plt.ylabel('Loss')
141     plt.legend()
142
143     plt.subplot(1, 2, 2)
144     plt.plot(range(1, 1+epochs), train_accuracy_values, label='Training Accuracy')
145     plt.plot(range(1, 1+epochs), test_accuracy_values, label='Test Accuracy')
146     plt.xlabel('Epochs')
147     plt.ylabel('Accuracy (%)')
148     plt.legend()
149
150     plt.show()

```